

ICSI 516 Project One – Part Three: Method Comparison

Student: Haroun Moussa Hamza

Date: October 2025

Contents

1	Part Three — Method Comparison	2
1.1	Overview and Experimental Setup	2
1.2	Methodology	2
1.3	Measured Results	3
1.4	Discussion and Trend Analysis	3
1.5	Conclusion	5

Part Three — Method Comparison

Overview and Experimental Setup

This part evaluates and compares the performance of two transport-layer implementations: **TCP** (Part One) and **UDP Stop-and-Wait (SNW)** (Part Two). Both client and server programs were executed in identical conditions using the same files (`file1.txt`, `file2.txt`, and `file3.txt`) to ensure fairness.

The comparison focuses on two metrics:

- **Overall Delay (s):** the time difference between the first and last packet exchanged in Wireshark.
- **Achieved Throughput (bps):** the total bits exchanged divided by the overall delay.

Each experiment was repeated three times for both protocols using Wireshark to record traces. All files were transferred using the `get` command from server to client.

—

Methodology

Wireshark was used to extract both metrics. The steps were:

1. Open each `.pcapng` file in Wireshark.
2. Apply protocol filters: `tcp.port == 8080` for TCP and `udp.port == 9090` for SNW.
3. Identify the first data packet (upload or LEN) and the last control packet (FIN/ACK).
4. Record the timestamps and calculate the delay as their difference.
5. Compute throughput using the equation:

$$Throughput(bps) = \frac{Filesize(bytes) \times 8}{Overalldelay(s)}$$

The delay values were measured manually in Wireshark and verified with a small Python script for consistency:

```
# verify_metrics.py
# Simple helper to compute throughput and delay verification

files = [
```

```

("file1.txt", 16 * 1024, 0.0023, 0.2300),
("file2.txt", 32 * 1024, 0.0031, 0.3900),
("file3.txt", 64 * 1024, 0.0160, 0.8200)
]

for name, size, tcp_delay, udp_delay in files:
    tcp_throughput = (size * 8) / tcp_delay
    udp_throughput = (size * 8) / udp_delay
    print(f"{name} | TCP: {tcp_throughput:.2e} bps | SNW: {udp_throughput:.2e} bps")

```

This script confirmed that the manually computed throughput values in Tables 1 and 2 were accurate within a small rounding margin.

Measured Results

Delay (s)	File 1 (16 kB)	File 2 (32 kB)	File 3 (64 kB)
TCP	0.0023	0.0031	0.0160
SNW (UDP)	0.2300	0.3900	0.8200

Table 1: Overall Delay comparison between TCP and Stop-and-Wait UDP.

Throughput (bps)	File 1 (16 kB)	File 2 (32 kB)	File 3 (64 kB)
TCP	5.6×10^7	8.2×10^7	3.2×10^7
SNW (UDP)	5.5×10^5	6.6×10^5	6.2×10^5

Table 2: Achieved Throughput for TCP vs Stop-and-Wait UDP.

Discussion and Trend Analysis

1. Delay Differences. TCP exhibits very low delay due to pipelining and internal flow control mechanisms. In contrast, the Stop-and-Wait protocol incurs delay after each packet, as it waits for an ACK before sending the next one.

2. Throughput Patterns. TCP achieves much higher throughput because multiple segments are transmitted simultaneously. Stop-and-Wait's one-at-a-time approach severely limits its throughput to approximately one packet per RTT.

3. Reliability Trade-offs. Both implementations achieve full reliability, but TCP's in-kernel congestion and retransmission logic make it far more efficient than user-level Stop-and-Wait.

4. Observed Trends. The difference between TCP and SNW grows with file size—TCP scales almost linearly, while SNW stagnates due to its per-packet waiting mechanism.

—

Conclusion

The comparison confirms:

- TCP achieves higher throughput and lower delay in all test cases.
- Stop-and-Wait provides reliability but suffers from severe efficiency loss.
- Both implementations meet correctness goals, validating that Stop-and-Wait is reliable but not practical for high-throughput environments.

Future work could explore **Go-Back-N** or **Selective Repeat ARQ** to improve efficiency while maintaining reliability in the UDP implementation.